

# Time Series Analysis

1. Imputation using interpolation
2. Decomposition into trend, seasonalit, residuals using statsmodels.api library
3. Exponential Smoothing: Simple Exp Smoothing(SES),Double Exp Smoothing(DES), Tripple Exponential Smoothing(TES)
4. Converting non stationary TS to Stationary TS
5. Types of Models: AR(Auto Regression), MA( Moving Average), ARIMA, SARIMA

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [36]: data=pd.read_excel(r"C:\Users\kanwar\Downloads\mobilesales.xlsx")
```

```
In [3]: data.head()
```

```
Out[3]:
```

	DATE	Sales
0	2001-01-01	6519.0
1	2001-02-01	6654.0
2	2001-03-01	7332.0
3	2001-04-01	7332.0
4	2001-05-01	8240.0

```
In [4]: data.tail()
```

Out[4]:

	DATE	Sales
212	2018-09-01	13838.0
213	2018-10-01	15351.0
214	2018-11-01	15615.0
215	2018-12-01	16941.0
216	2019-01-01	12160.0

In [5]: `data.dtypes`

Out[5]: DATE        datetime64[ns]  
Sales            float64  
dtype: object


In [6]: `data.describe()`

Out[6]:

	DATE	Sales
count	217	198.000000
mean	2009-12-31 04:25:26.267281152	10842.757576
min	2001-01-01 00:00:00	108.000000
25%	2005-07-01 00:00:00	8856.000000
50%	2010-01-01 00:00:00	10745.000000
75%	2014-07-01 00:00:00	12399.750000
max	2019-01-01 00:00:00	23740.000000
std	NaN	2605.006293


In [37]: `data.set_index('DATE', inplace=True)`

```
In [38]: data.plot();
```




## Imputation

```
In [39]: data['Sales'].fillna(0).plot(label='imputed')  
data['Sales'].plot(label='original')  
plt.legend();
```



```
In [40]: data['Sales'].fillna(data['Sales'].mean()).plot(label="imputed")
data['Sales'].plot(label="original")


plt.legend()
plt.show()
```



```
In [41]: # linear interpolation


data.Sales.interpolate(method = 'linear').plot(label='imputed')
data.Sales.plot(label='original')

plt.legend();
```




```
In [42]: data.Sales = data.Sales.interpolate(method = 'linear')
data.Sales.plot()
```

```
Out[42]: <Axes: xlabel='DATE'>
```




```
In [13]: data.Sales.hist(bins=100).plot()
```

```
Out[13]: []
```




```
In [43]: data.Sales.clip(upper=data.Sales.quantile(0.95),lower=data.Sales.quantile(0.05)).plot()  
data['Sales'].plot(label="original")  
plt.legend();
```






```
In [44]: data.Sales.clip(upper=data.Sales.quantile(0.95),lower=data.Sales.quantile(0.05)).plot();
```




```
In [45]: data.Sales=data.Sales.clip(upper=data.Sales.quantile(0.95),lower=data.Sales.quantile(0.05))
```

```
In [46]: data.Sales.rolling(12).mean().plot(label='MA') ## Moving Average (MA) 12 months  
data.Sales.plot(label='original')  
plt.legend();
```




```
In [47]: import statsmodels.api as sm  
model=sm.tsa.seasonal_decompose(data.Sales)
```


```
In [24]: model.plot();
```



```
In [48]: plt.subplot(2,1,1)
model.trend.plot()
plt.subplot(2,1,2)
data.Sales.rolling(12,center=True).mean().plot();
```



```
In [49]: plt.subplot(2,1,1)
model.seasonal.plot()
plt.subplot(2,1,2)
data.Sales.hist(bins=100);
```



```
In [39]: model.resid.mean()
```

```
Out[39]: np.float64(-2.547773692809866)
```

```
In [53]: train_max_date = data.index[-12]
```

```
train_x = data.loc[data.index < data.index[-12]].copy()  
test_x = data.loc[data.index >= data.index[-12]].copy()
```

```
In [41]: train_x.shape, test_x.shape
```

```
Out[41]: ((205, 1), (12, 1))
```

```
In [42]: test_x
```

Out[42]:

**Sales**

DATE	
2018-02-01	11852.0
2018-03-01	14123.0
2018-04-01	13360.0
2018-05-01	15164.6
2018-06-01	15164.6
2018-07-01	14080.0
2018-08-01	15164.6
2018-09-01	13838.0
2018-10-01	15164.6
2018-11-01	15164.6
2018-12-01	15164.6
2019-01-01	12160.0

```
In [50]: from sklearn.metrics import (
    mean_squared_error as mse,
    mean_absolute_error as mae,
    mean_absolute_percentage_error as mape
)

# Creating a function to print values of all these metrics.
def performance(actual, predicted):
    print('MAE :', round(mae(actual, predicted), 3))
    print('RMSE :', round(mse(actual, predicted)**0.5, 3))
    print('MAPE:', round(mape(actual, predicted), 3))
```

```
In [44]: train_x['Sales'].mean()
```

```
Out[44]: np.float64(10659.16)
```


```
In [45]: # Mean model
test_x['pred'] = train_x['Sales'].mean()
test_x.head()
```

```
Out[45]:
```

	Sales	pred
DATE		
2018-02-01	11852.0	10659.16
2018-03-01	14123.0	10659.16
2018-04-01	13360.0	10659.16
2018-05-01	15164.6	10659.16
2018-06-01	15164.6	10659.16

```
In [46]: test_x.plot(style='-o');
```





```
In [47]: performance(test_x['Sales'], test_x['pred'])
```


```
MAE : 3540.89  
RMSE : 3725.69  
MAPE: 0.244
```

```
In [50]: ## Naive Model: Prediction is the last value of the train data
```

```
In [54]: test_x['pred'] = train_x['Sales'][-1]
```

```
C:\Users\kanwar\AppData\Local\Temp\ipykernel_28896\4166961390.py:1: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`  
test_x['pred'] = train_x['Sales'][-1]
```

```
In [55]: test_x.plot(style='-o');
```



```
In [51]: performance(test_x['Sales'], test_x['pred'])
```

```
MAE : 3195.05  
RMSE : 3398.707  
MAPE: 0.219
```

```
In [56]: # Seasonal Naive
pred = []


for i in test_x.index:
    pred.append(train_x.loc[i - pd.DateOffset(years=1)][ 'Sales' ])

test_x['pred'] = pred
test_x.head()
```

```
Out[56]:
```

	Sales	pred
DATE		
2018-02-01	11852.0	11891.0
2018-03-01	14123.0	13921.0
2018-04-01	13360.0	14453.5
2018-05-01	15164.6	14986.0
2018-06-01	15164.6	15164.6


```
In [57]: test_x.plot(style='-o');
```



```
In [54]: performance(test_x['Sales'], test_x['pred'])
```

```
MAE : 660.367  
RMSE : 905.569  
MAPE: 0.047
```

```
In [55]: plt.plot(train_x['Sales'][-48:], label='train')  
plt.plot(test_x)  
plt.plot();
```



```
In [56]: ## Drift Method: Naive+Trend

# Get the slope
yt = train_x['Sales'][-1]


m = (yt - train_x['Sales'][0]) / len(train_x)

h = np.arange(0, 12)

test_x['pred'] = yt + h*m
```

```
C:\Users\kanwar\AppData\Local\Temp\ipykernel_18508\1029698906.py:4: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
yt = train_x['Sales'][-1]
C:\Users\kanwar\AppData\Local\Temp\ipykernel_18508\1029698906.py:6: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
m = (yt - train_x['Sales'][0]) / len(train_x)
```


```
In [57]: test_x.plot(style='-o');
```



```
In [58]: performance(test_x['Sales'], test_x['pred'])
```


MAE : 3095.857  
RMSE : 3301.099  
MAPE: 0.212

```
In [59]: train_x.Sales.plot();  
train_x.Sales.rolling(12).mean().plot();
```



```
In [60]: pred_ = train_x.Sales.copy().values  
  
for i in range(12):  
    pred_ = np.append(pred_, pred_[-6:].mean()) # MA(6) last 6 months avg.  
  
test_x['pred'] = pred_[-12:] # predictions
```

```
In [61]: test_x.plot(style='-o');
```



```
In [62]: performance(test_x.Sales, test_x.pred);
```

MAE : 1247.799  
RMSE : 1364.818  
MAPE: 0.088

## Exponential Smoothing Model:

1. Simple Exp Smoothing model
2. Double Exp Smoothing model



### 3. Tripple Exp Smoothing Model

```
In [ ]: ## SESM : Naive+Rolling mean with weightage which reduce exponentially
```

```
In [63]: model = sm.tsa.SimpleExpSmoothing(data.Sales).fit(smoothing_level=0.2) # alpha  
pred = model.forecast(12) # pred for next 12 months.
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```


```
    self._init_dates(dates, freq)
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\pandas\util\_decorators.py:213: EstimationWarning: Model has no free parameters to estimate. Set optimized=False to suppress this warning
```

```
    return func(*args, **kwargs)
```

```
In [64]: data.Sales[-120:].plot(label='original', style='-o')  
  
pred.plot(label='ses_pred', style='-o')  
plt.legend()
```


```
Out[64]: <matplotlib.legend.Legend at 0x169038e5f90>
```



```
In [14]: # just by adding trend -> DES
model = sm.tsa.ExponentialSmoothing(data.Sales, trend="add").fit(smoothing_trend=0.02);
model.fittedvalues.plot(label="fitted")
data.Sales.plot(label="original")

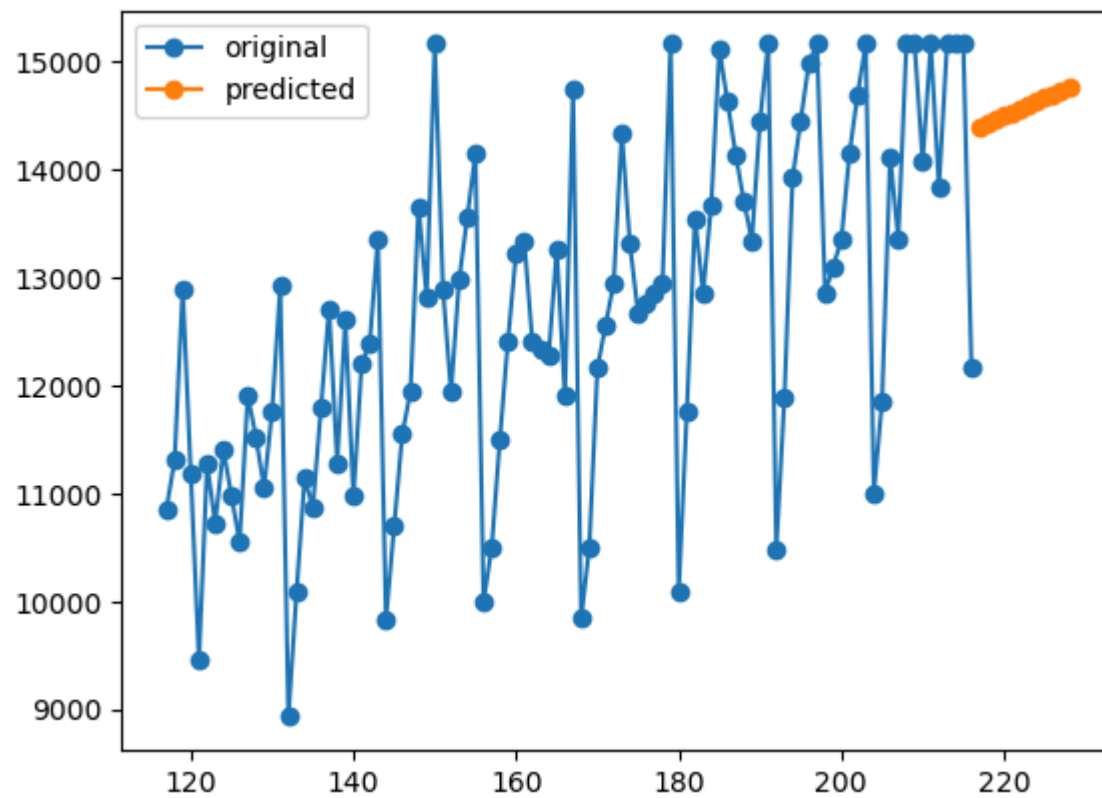
plt.legend();
```

C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\holtwinters\model.py:918: ConvergenceWarning: Optimization failed to converge. Check mle\_retvals.  
warnings.warn(



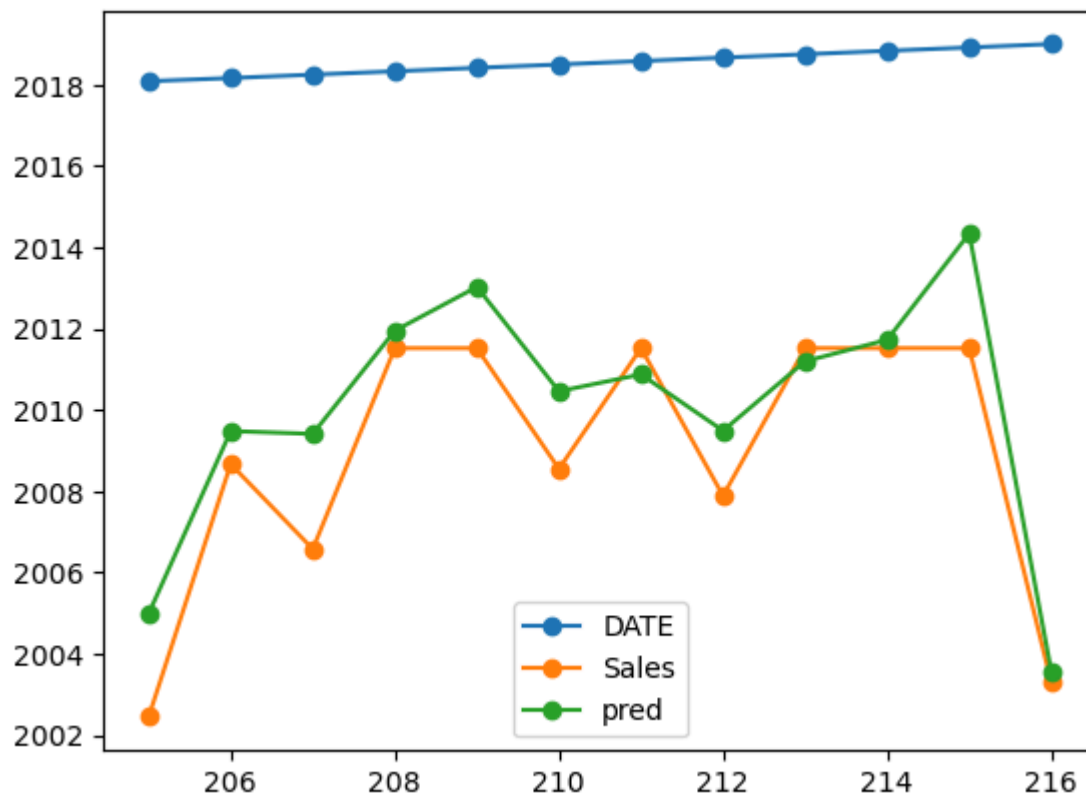
```
In [16]: pred = model.forecast(12)
data.Sales.tail(100).plot(label="original", style='-o')
pred.plot(label="predicted", style='-o')
plt.legend()
```

```
Out[16]: <matplotlib.legend.Legend at 0x16800e64550>
```



```
In [52]: test_x.plot(style='-o')  
  
performance(test_x['Sales'], test_x['pred'])
```

MAE : 483.93  
RMSE : 599.863  
MAPE: 0.035




```
In [59]: model = sm.tsa.ExponentialSmoothing(train_x.Sales, trend="add", seasonal = "add").fit();
```

C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.  
 self.\_init\_dates(dates, freq)

```
In [61]: pred = model.forecast(12)
train_x.Sales.tail(100).plot(label="original", style='-o')
pred.plot(label="predicted", style='-o')
plt.legend()
```


Out[61]: <matplotlib.legend.Legend at 0x1680801b9d0>



```
In [62]: test_x['pred']=pred.values
```

```
In [63]: test_x.plot(style='-o')  
  
performance(test_x['Sales'], test_x['pred'])
```

MAE : 473.509  
RMSE : 565.711  
MAPE: 0.033

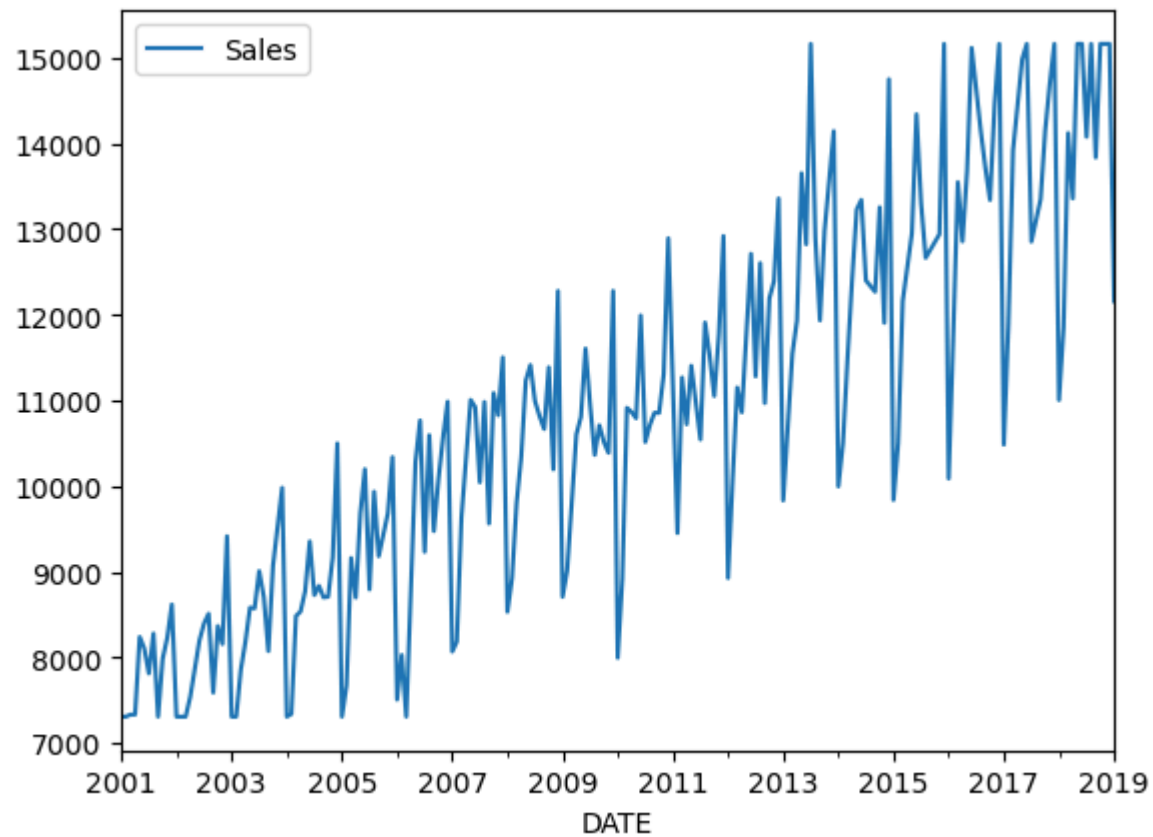


```
In [64]: ## check if the time series is stationary:

def adf_test(data, significance_level=0.05):
    pvalue = sm.tsa.stattools.adfuller(data)[1]

    if pvalue <= significance_level:
        print('Sequence is stationary')
    else:
        print('Sequence is not stationary')
```

```
In [66]: data.plot();
```



```
In [67]: adf_test(data.Sales)
```


Sequence is not stationary

```
In [69]: ### Convert Non Stationary Time series to Stationary Time Series
```

```
model=sm.tsa.seasonal_decompose(data.Sales)
```


```
In [71]: model.plot();
```





```
In [75]: model.resid.plot()
```

```
Out[75]: <Axes: xlabel='DATE'>
```



## Conversion of Non Stationary to Stationary using Residual

Method1: Decomposition

$$y(t) = b(t) + s(t) + e(t)$$

Remove trends from  $e(t)$  which has no trend and seasonality which makes it a stationary time series.


$$e(t) = y(t) - b(t) - s(t)$$

Method2: Differentiation

```
In [77]: adf_test(model.resid.dropna())
```

Sequence is stationary

```
In [82]: ## Differentiation  
data.Sales.diff(1).plot();
```



```
In [84]: adf_test(data.Sales.diff(1).dropna())
```

Sequence is stationary

```
In [85]: ## If this is non stationary then we can do double differentiation i.e make it more smoother
```

## Auto correlation

```
In [89]: data.Sales.shift(1).head()
```

```
Out[89]: DATE
2001-01-01      NaN
2001-02-01    7307.8
2001-03-01    7307.8
2001-04-01    7332.0
2001-05-01    7332.0
Name: Sales, dtype: float64
```

```
In [91]: ## check auto correlation with lag=12
lag=12
np.corrcoef(data.Sales[lag:],data.Sales.shift(lag)[lag:])[0][1]
```

```
Out[91]: np.float64(0.9419098493883264)
```


```
In [92]: ## check auto correlation with lag=12
lag=6
np.corrcoef(data.Sales[lag:],data.Sales.shift(lag)[lag:])[0][1]
```

```
Out[92]: np.float64(0.7966090507569205)
```

```
In [93]: ## There is high correlation with 12 lag i.e because there is seasonality of 12 months
```


```
In [94]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [95]: plot_acf(data.Sales);
```




```
In [96]: data.Sales.diff(1).plot()
```

```
Out[96]: <Axes: xlabel='DATE'>
```



```
In [99]: ## Differentiation removes the trend from the data  
plot_acf(data.Sales.diff().dropna());
```




**Highest correlation is visible with lag value=12**

Note here the correlation is also negative because the differentiation values are negative as well which is unlike the original data

In [100... `## Lecture:4`

**ARIMA:**

In [101... `train_x.plot();`



```
In [102... train_x_st = train_x.copy()


train_x_st.Sales = train_x_st.Sales.diff(1)
train_x_st.dropna(inplace=True)
```

```
In [103... adf_test(train_x_st.Sales)
```

Sequence is stationary

```
In [104... train_x_st.Sales.plot();
```





```
In [105... from statsmodels.tsa.statespace.sarimax import SARIMAX # supermodel
```

```
In [106... # order = (p, d, q)
model = SARIMAX(train_x_st.Sales, order = (3,0,0) ) # AR(3)
model = model.fit(dispatch = False)
test_x['pred'] = model.forecast(steps=12)
```

C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

```
In [107... test_x.head(10)
```

```
Out[107...
```

	Sales	pred
DATE		
2018-02-01	11852.0	1806.871390
2018-03-01	14123.0	512.621534
2018-04-01	13360.0	-6.529820
2018-05-01	15164.6	-582.530925
2018-06-01	15164.6	183.503533
2018-07-01	14080.0	122.694532
2018-08-01	15164.6	-1.070891
2018-09-01	13838.0	-84.884694
2018-10-01	15164.6	16.220618
2018-11-01	15164.6	23.212393

## Why negative sales??

Because the model is trained with differences/differentiated data, hence post processing is required to check performance to compare the prediction and actual value

```
In [109... # order = (p, d, q)
model = SARIMAX(train_x_st.Sales, order = (3,0,0) ) # AR(3)
model = model.fit(dispatch = False)
test_x['pred'] = model.forecast(steps=12)


# post processing
test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\Users\kanwar\AppData\Local\Temp\ipykernel_28896\2208145763.py:7: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()
```

```
In [110... test_x.plot(style='-o')

performance(test_x['Sales'], test_x['pred'])
```

```
MAE : 1489.515
RMSE : 1677.498
MAPE: 0.102
```



```
In [111... # order = (p, d, q)
model = SARIMAX(train_x_st.Sales, order = (12,0,0) ) # AR(12)
model = model.fit(dispatch = False)
test_x['pred'] = model.forecast(steps=12)

# post processing
test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()

test_x.plot(style='-o')

performance(test_x['Sales'], test_x['pred'])
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```


```
C:\Users\kanwar\AppData\Local\Temp\ipykernel_28896\2927954615.py:7: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
```

```
test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()
```

MAE : 747.487

RMSE : 897.502

MAPE: 0.053



In [112...

```
# order = (p, d, q)
model = SARIMAX(train_x_st.Sales, order = (0,0,10) ) # MA(10)
model = model.fit(dispatch = False)
test_x['pred'] = model.forecast(steps=12)

# post processing
test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()

test_x.plot(style='-o')

performance(test_x['Sales'], test_x['pred'])
```

C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.

```
self._init_dates(dates, freq)
```

C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle\_retvals

```
warnings.warn("Maximum Likelihood optimization failed to "
```


C:\Users\kanwar\AppData\Local\Temp\ipykernel\_28896\3566005573.py:7: FutureWarning: Series.\_\_getitem\_\_ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

```
test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()
```

MAE : 1095.166

RMSE : 1369.994

MAPE: 0.075



```
In [113... # order = (p, d, q)
model = SARIMAX(train_x_st.Sales, order = (12,0,12) ) # ARMA(12,12)
model = model.fit(dispatch = False)
test_x['pred'] = model.forecast(steps=12)

# post processing
test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()


test_x.plot(style='-o')

performance(test_x['Sales'], test_x['pred'])
```


```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters.
    warn('Non-stationary starting autoregressive parameters')
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
    warn('Non-invertible starting MA parameters found.')
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to ")
C:\Users\kanwar\AppData\Local\Temp\ipykernel_28896\2884599505.py:7: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    test_x['pred'] = train_x.Sales[-1] + test_x['pred'].cumsum()

MAE : 619.802
RMSE : 713.085
MAPE: 0.045
```





In [115... `train_x.plot();`



In [116... `## ARIMA: I for integration hence it handles post processing and preprocessing, AR+I+MA`

In [120... `# order = (p, d, q)`  
`model = SARIMAX(train_x.Sales, order = (5,1,3) ) # ARIMA(5, 1, 3) #5 previous values , 3 previous errors and 1 differentiation`  
`model = model.fit(dispatch = False)`  
`test_x['pred'] = model.forecast(steps=12)`  
  
`test_x.plot(style='-o')`  
`performance(test_x['Sales'], test_x['pred'])`

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```


```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```

```
MAE : 1098.554
```

```
RMSE : 1198.155
```

```
MAPE: 0.077
```



In [121...

```
# (p, d, q)
model = SARIMAX(train_x.Sales, order = (15, 1, 12)) # ARIMA(5, 1, 5)
model = model.fit(dispatch = False)
test_x['pred'] = model.forecast(steps=12)
```

```
test_x.plot(style='-o')
performance(test_x['Sales'], test_x['pred'])
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
    self._init_dates(dates, freq)
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
    self._init_dates(dates, freq)
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters.
```

```
    warn('Non-stationary starting autoregressive parameters')
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
```

```
    warn('Non-invertible starting MA parameters found.')
```


```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
    warnings.warn("Maximum Likelihood optimization failed to "
```

```
MAE : 475.359
```

```
RMSE : 617.215
```

```
MAPE: 0.034
```



```
In [122... ## SARIMA: Seasonal ARIMA, seasonality+AR+I+MA
```

```
In [123... # order = (p, d, q), seasonal_order= P,D,Q,S---12 is seasonality, 12 months seasonality
model = SARIMAX(train_x.Sales, order = (5,1,3), seasonal_order=(1,1,1,12)) # SARIMA(5, 1, 3)
model = model.fit(dispatch = False)
test_x['pred'] = model.forecast(steps=12)

test_x.plot(style='-o')
performance(test_x['Sales'], test_x['pred'])
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\kanwar\anaconda3\Lib\site-packages\statsmodels\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

MAE : 426.811

RMSE : 508.667

MAPE: 0.031

